# Artificial Neural Network For Beginners

Admond Lee

# Little Bit About Me

## Admond Lee

Physics in NTU

Data Scientist

# Data Science Writer

# Data Science Consulting

- Helping digital marketing agencies and companies solve problems using data
- Data Science Consulting
  - ➢ Web Analytics
  - ➢ Attribution Modelling
  - ➢ Deep Learning / ML
  - ➢ Advanced Social Analytics

# Data Science Communicator


# What We'll Cover Today

**What is Artificial Neural Network?**

**Build a simple ANN model (for real)**

# What is Artificial Neural Network?

# Fundamental Building Blocks – Neurons

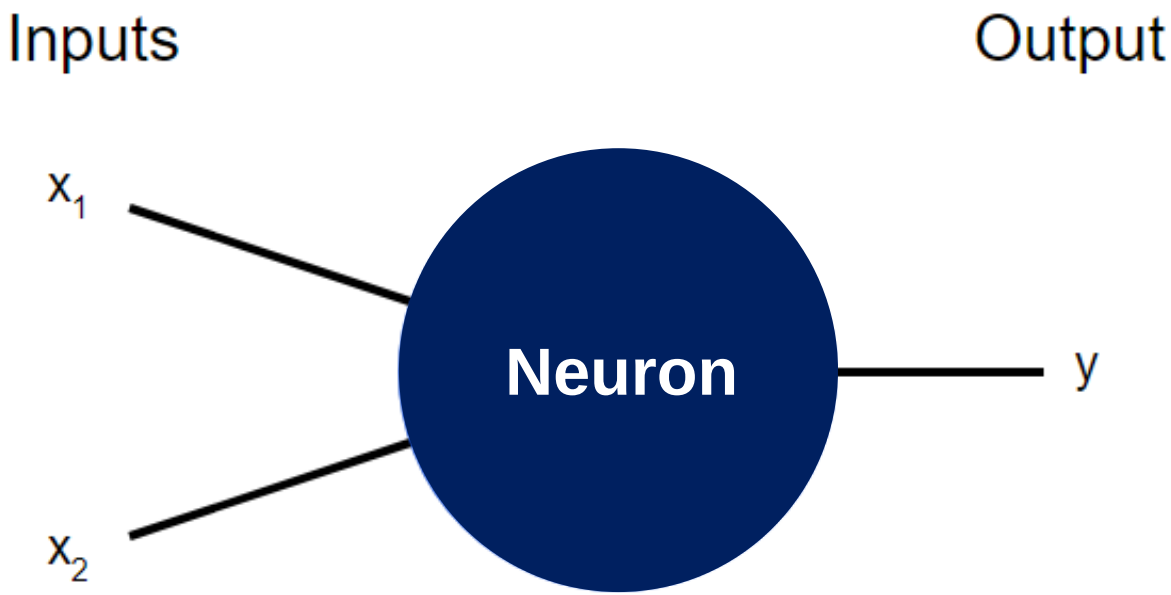


A neuron takes inputs, does some math with them,
and produces one output

# Fundamental Building Blocks – Neurons

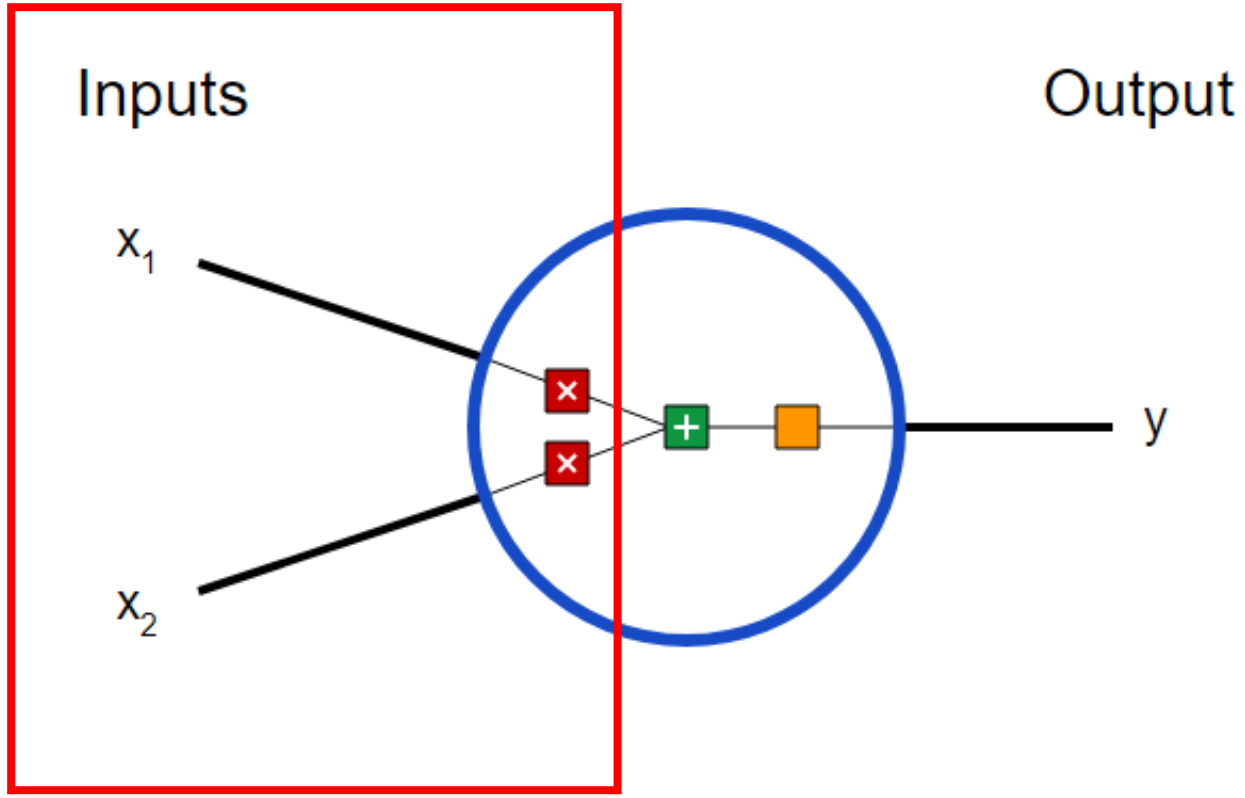


$x_1 \rightarrow x_1 * w_1$ weight

$x_2 \rightarrow x_2 * w_2$

The higher the weight:

- The faster it will trigger the activation function
- The more impact it will have on the network

# Fundamental Building Blocks – Neurons

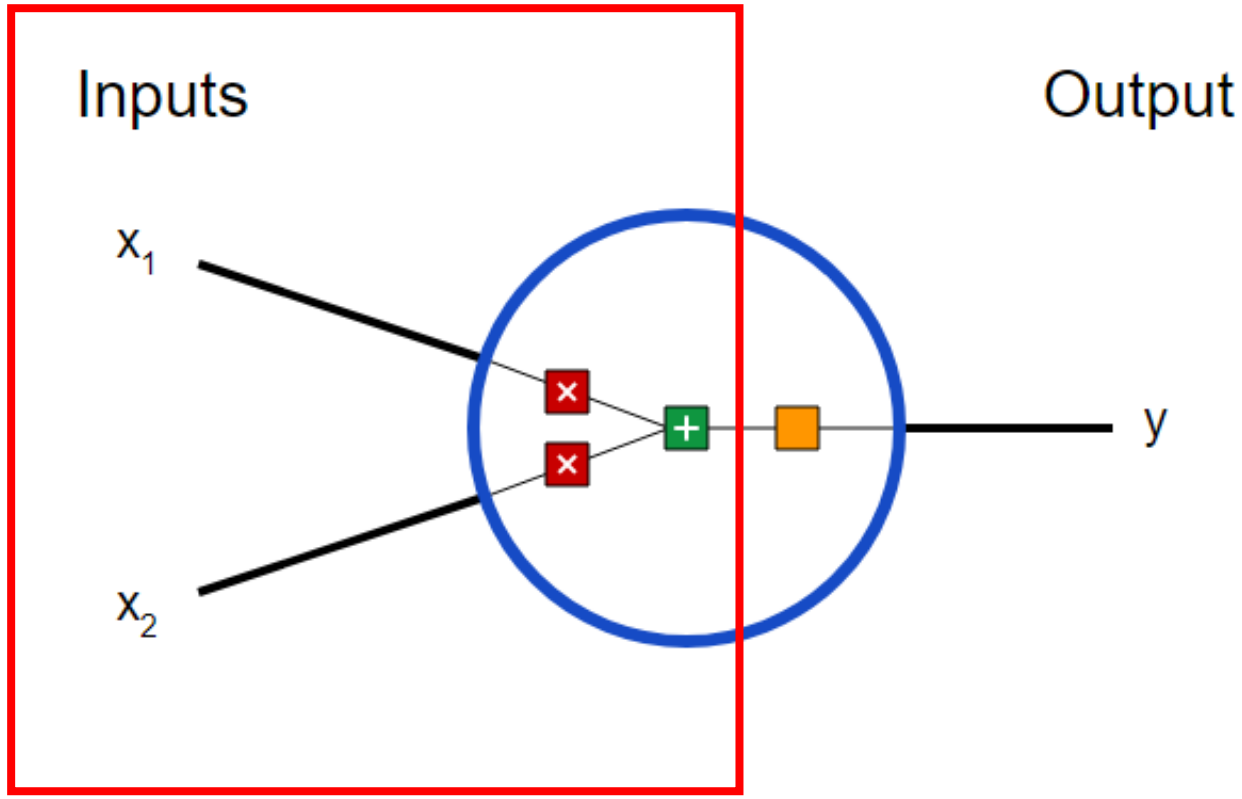


$$(x_1 * w_1) + (x_2 * w_2) + b$$ bias

Bias is used to delay the triggering of the activation function

# Fundamental Building Blocks – Neurons

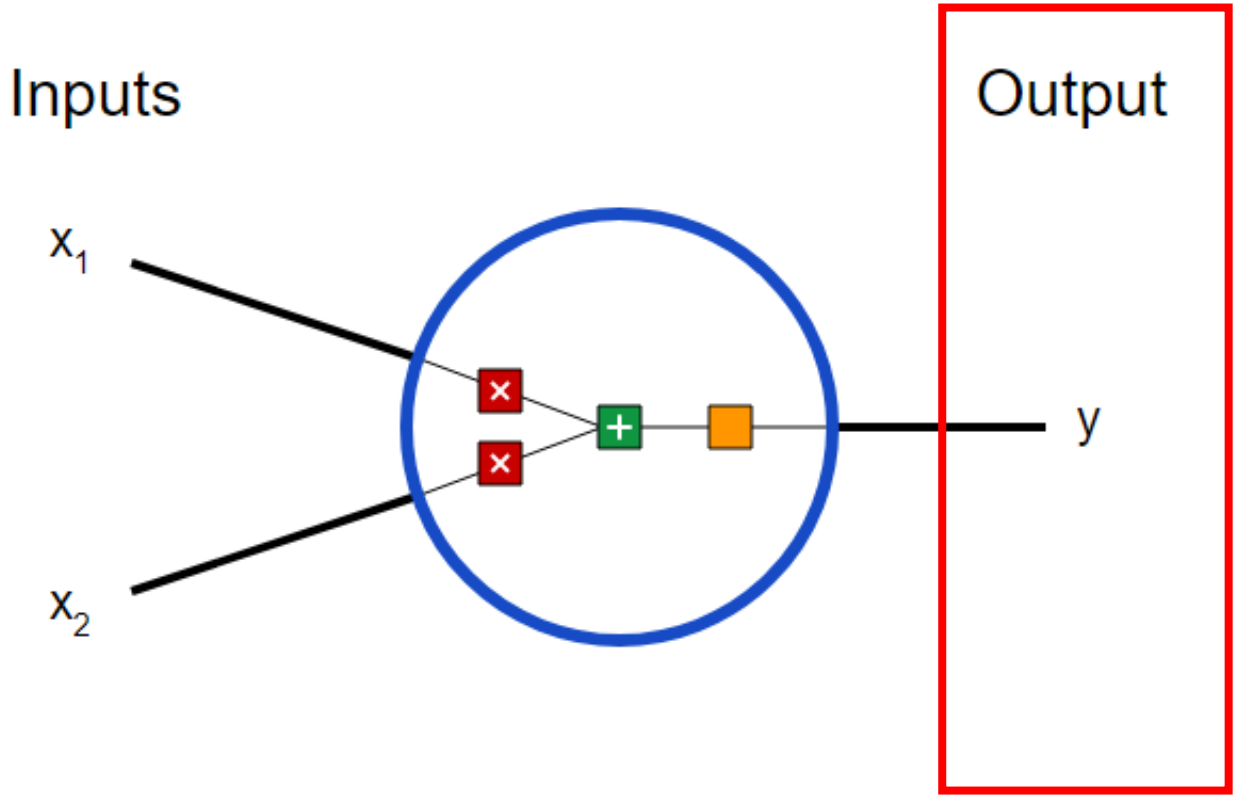


Activation function

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$

# Activation Function

- An activation function is used to turn an unbounded input into an output that has a nice, predictable form
- It will decide whether a neuron should fire or not

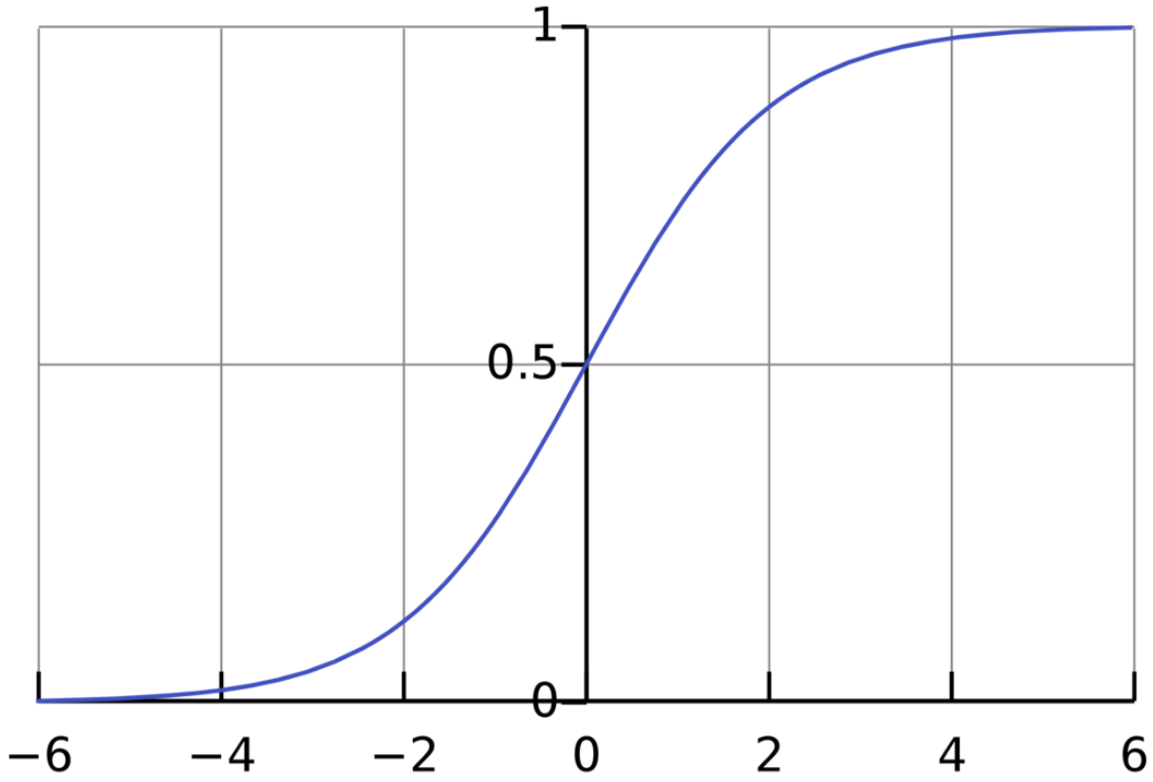


Sigmoid function

# Neural Network

- A neural network is nothing more than a bunch of neurons connected together
- A neural network can have any number of layers with any number of neurons in those layers

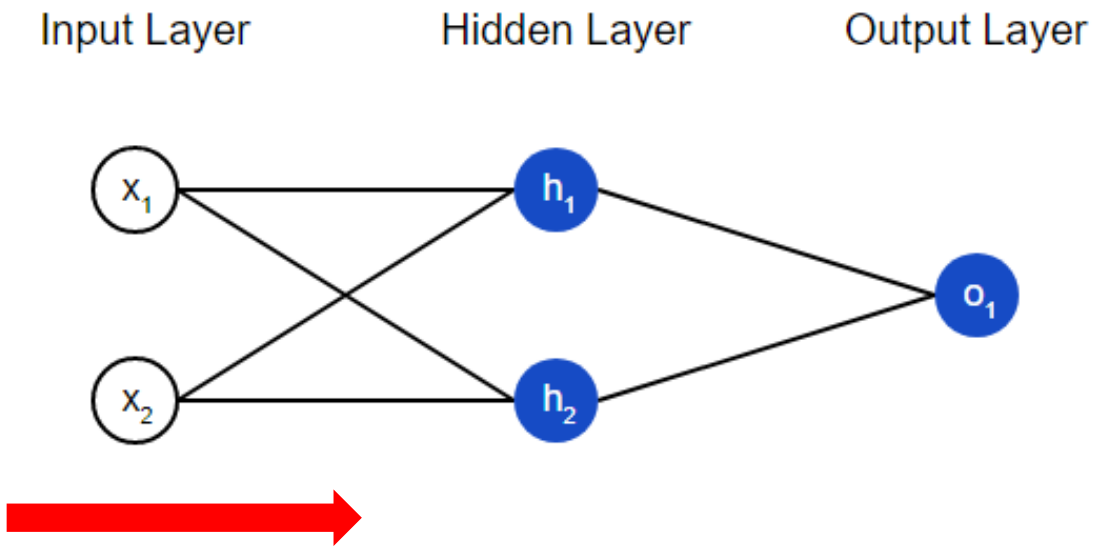


Feedforward Neural Network

# Training a Neural Network

- Minimize the error (or loss) from predictions
- Lower loss → Better predictions
- Mean Squared Error (MSE) → $\text{MSE} = \frac{1}{n}\sum_{i=1}^{n}(y_{true} - y_{pred})^2$ **(Loss function)**

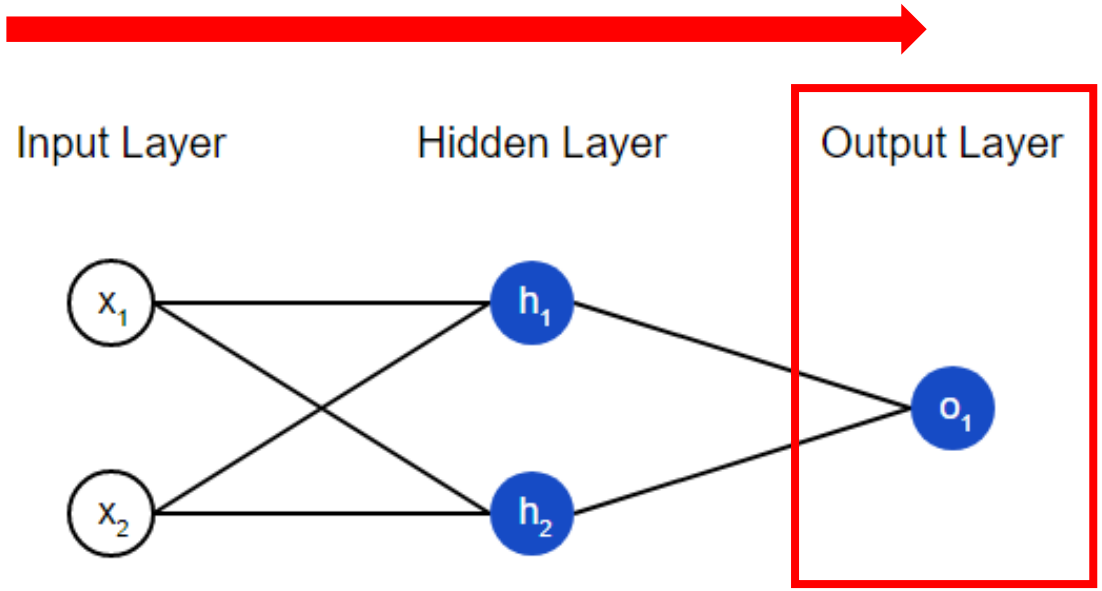


- Compare the actual and predicted values
- Compute loss function to find the error

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$

# Backpropagation

- AIM – Update the network's weights and biases to minimize the error from predictions
- Lower loss → Better predictions
- Mean Squared Error (MSE) → $\text{MSE} = \frac{1}{n}\sum_{i=1}^{n}(y_{true} - y_{pred})^2$ **(Loss function)**

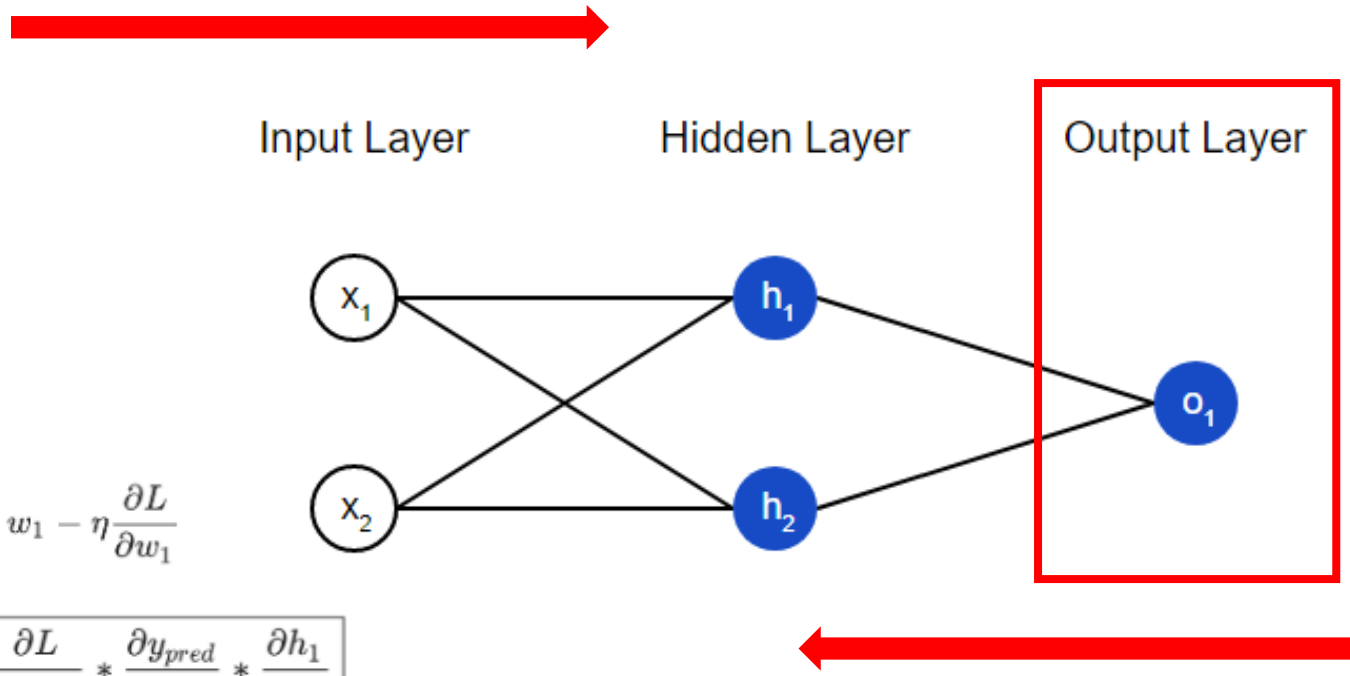


$$w_1 \leftarrow w_1 - \eta \frac{\partial L}{\partial w_1}$$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial y_{pred}} * \frac{\partial y_{pred}}{\partial h_1} * \frac{\partial h_1}{\partial w_1}$$

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$

- How to know which weights to be updated
  - ➢ Partial derivative
- This system of calculating partial derivatives by working backwards is known as backpropagation

# Example – Voice Recognition

**Task**

- Learn to discriminate between two different voices saying “Hello”

**Data**

- Sources
  - Steve
  - David
- Format
  - Frequency distribution (60 bins)

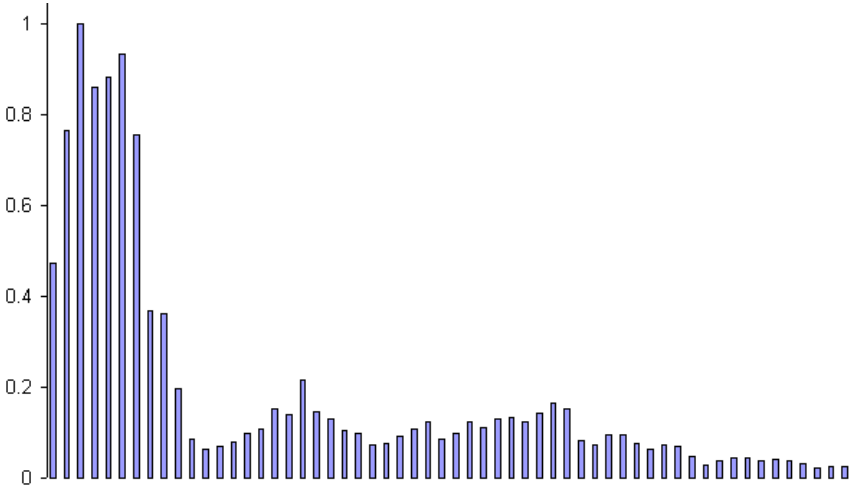


# Network architecture

- Feedforward network
  - 60 input (one for each frequency bin)
  - 6 hidden
  - 2 output (0-1 for “Steve”, 1-0 for “David”)

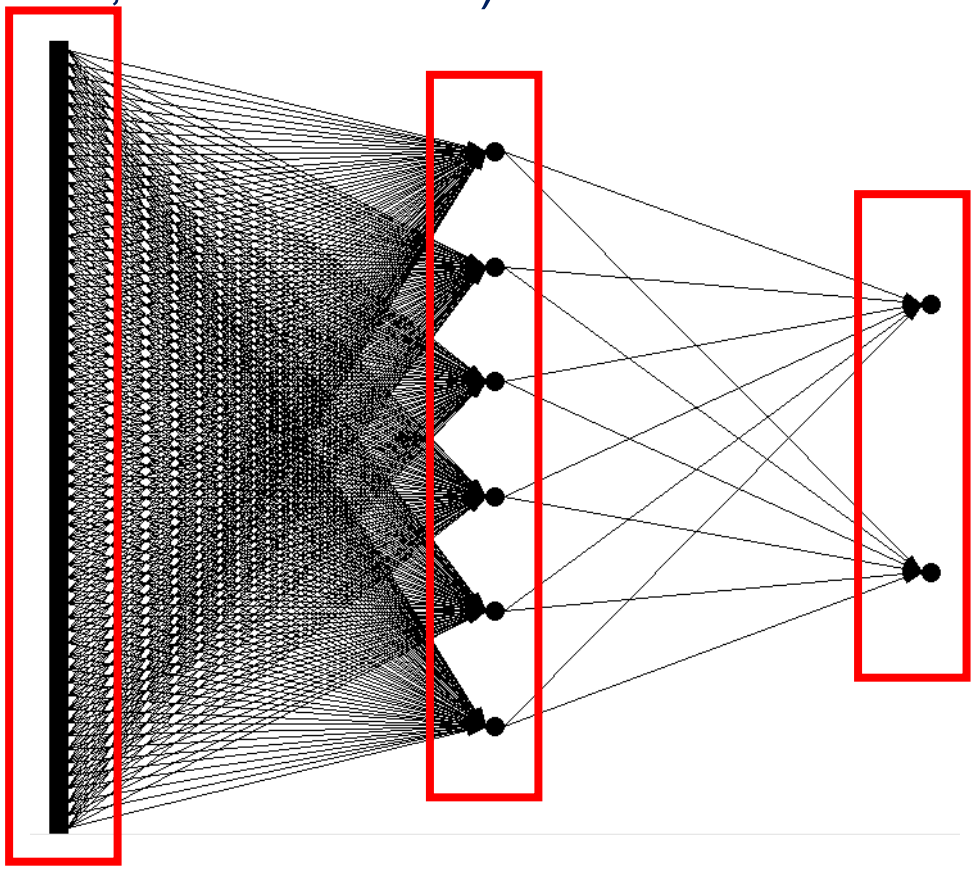

# Presenting the data

Steve

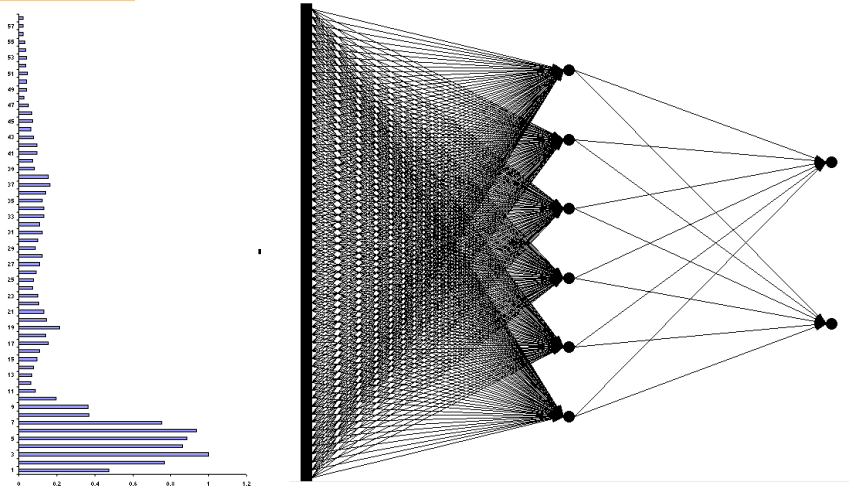

David

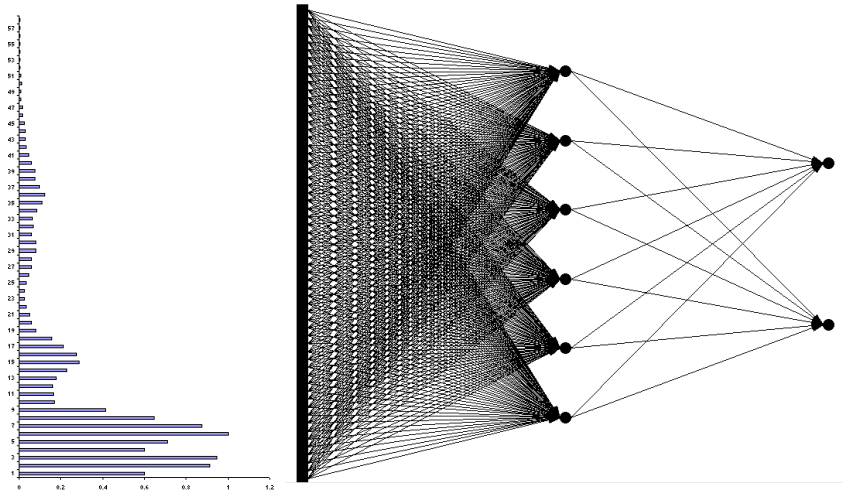

# Presenting the data (untrained network)

Steve

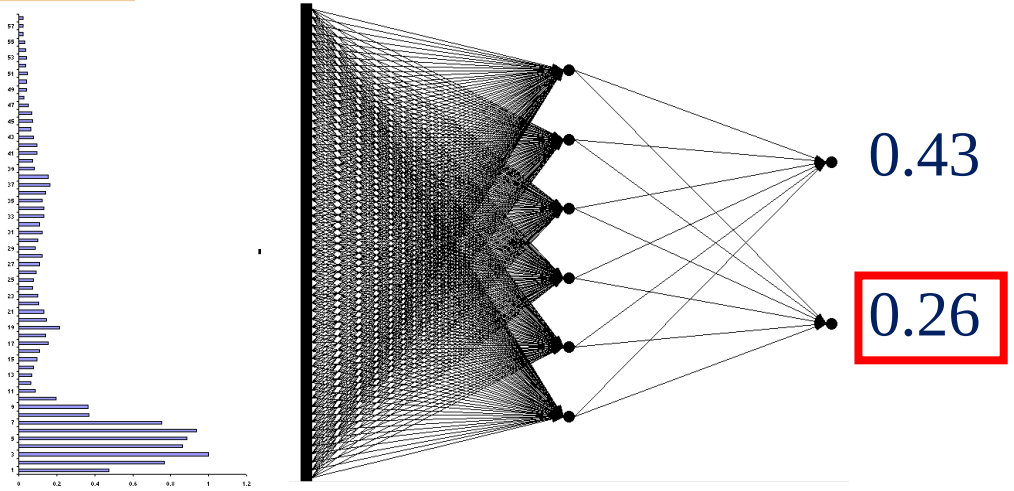


David

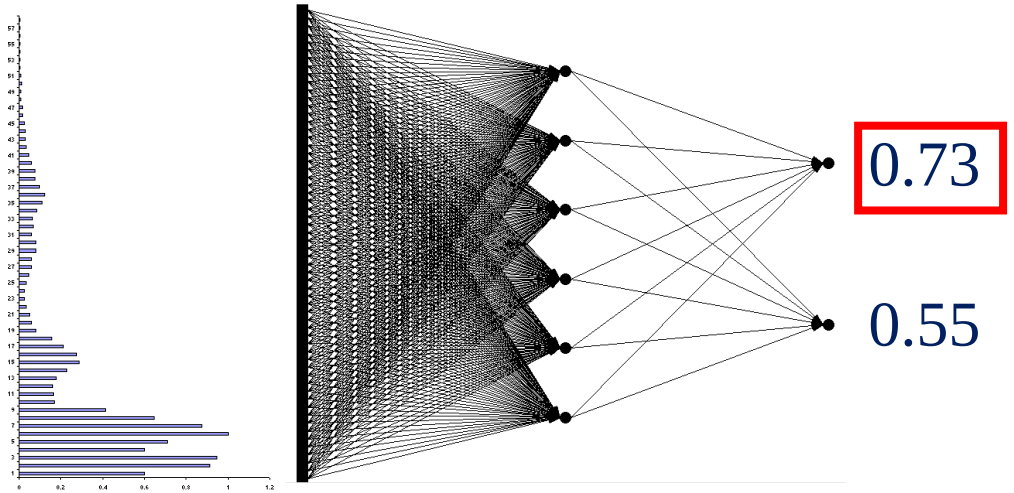

# Calculate error

Steve

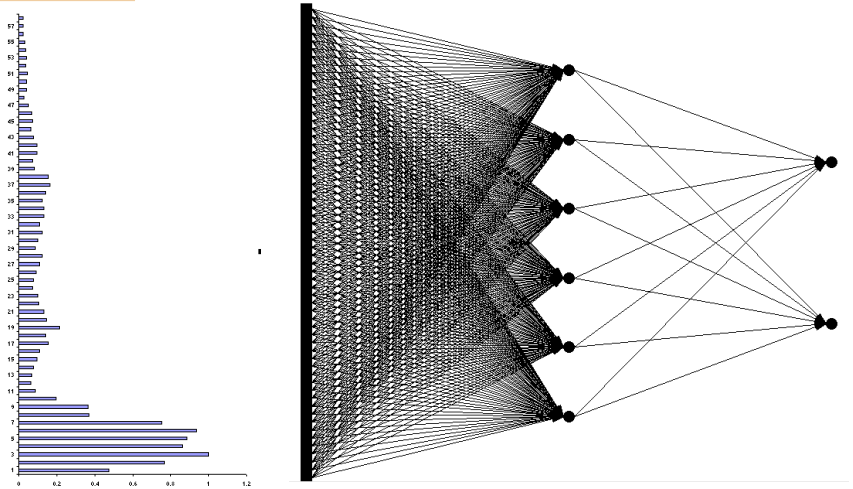

$0.43 - 0 = 0.43$

$0.26 - 1 = 0.74$

David

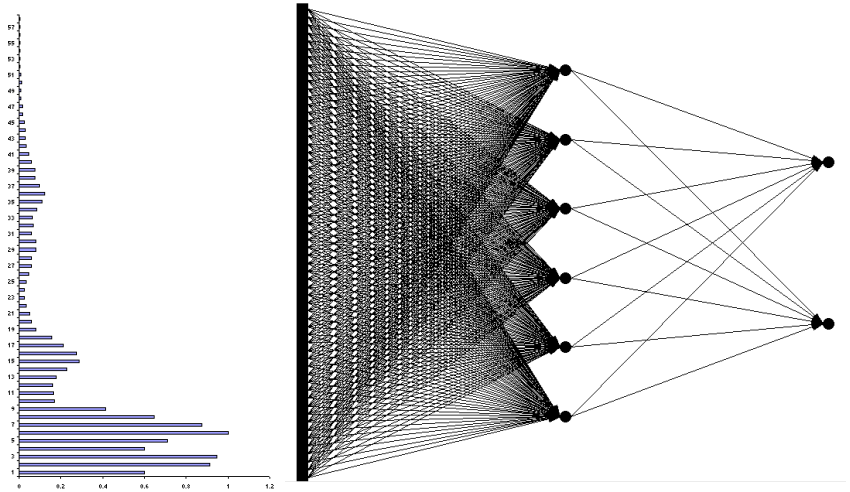

$0.73 - 1 = 0.27$

$0.55 - 0 = 0.55$

# Backpropagate error and adjust weights

Steve

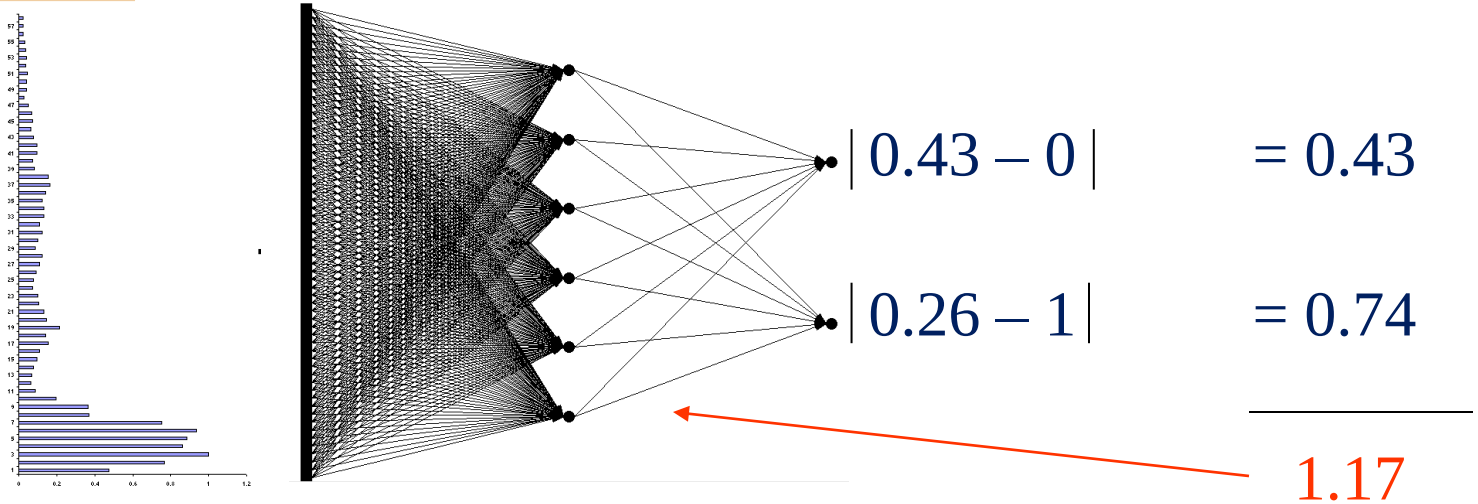

$$|0.43 - 0| = 0.43$$

$$|0.26 - 1| = 0.74$$

$$1.17$$

David

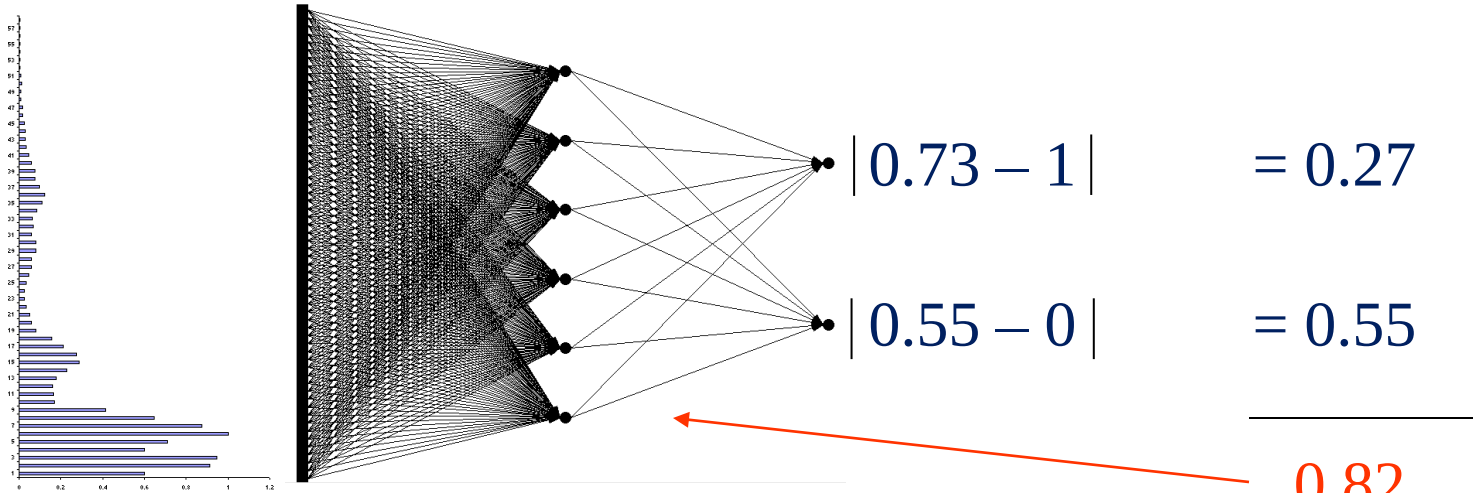

$$|0.73 - 1| = 0.27$$

$$|0.55 - 0| = 0.55$$

$$0.82$$

# Example – Voice Recognition

- Repeat process (sweep) for all training pairs
  - Present data
  - Calculate error
  - Backpropagate error
  - Adjust weights

- Repeat process multiple times

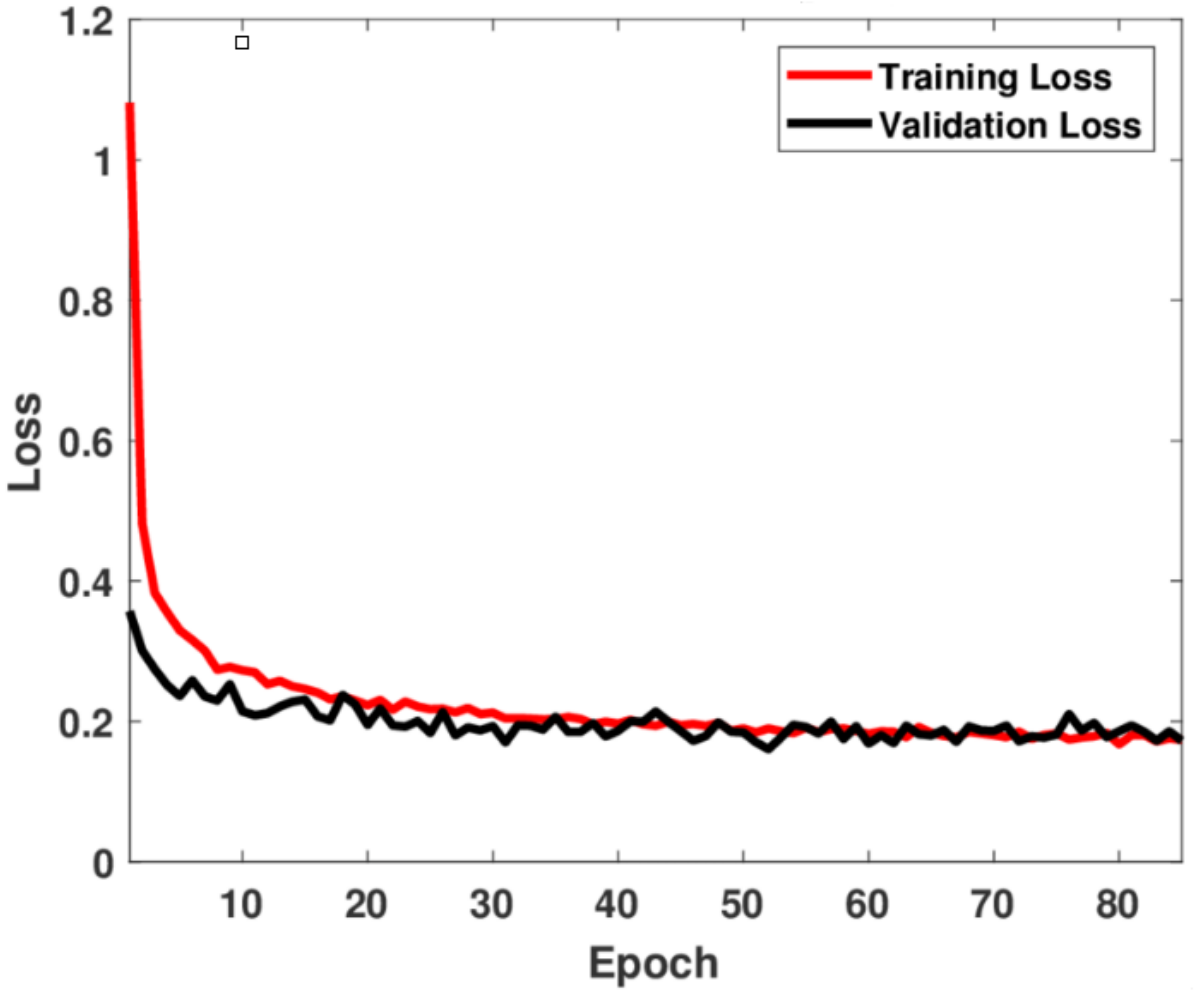

# Presenting the data (trained network)

Steve

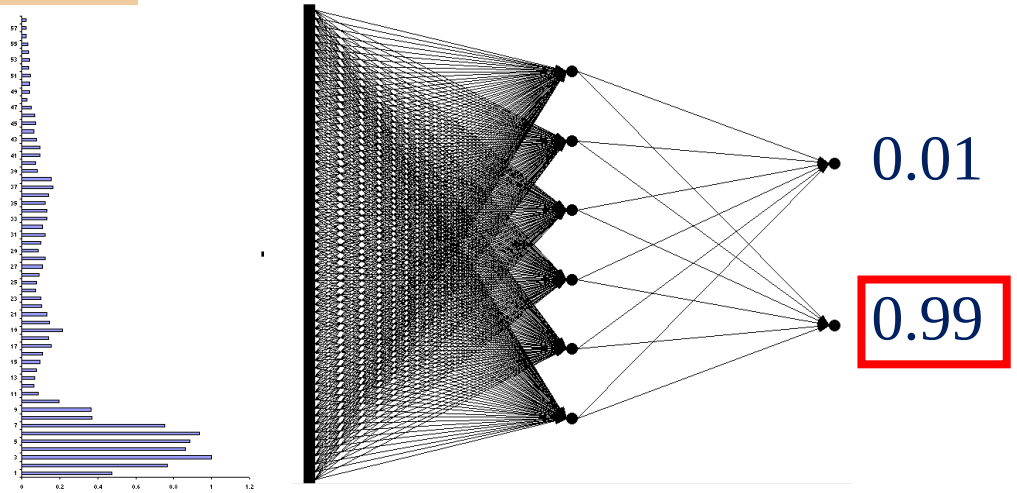


David

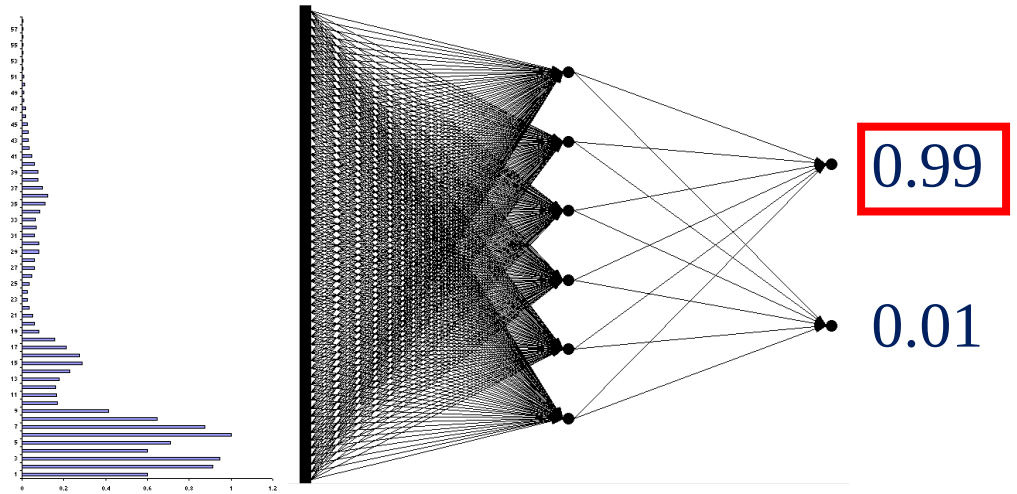

# Neural Network Demo

http://playground.tensorflow.org

Keras

# Introduction to Keras

# What is Keras?

- **Keras** is a high-level deep learning API written in Python that can run on top of the deep learning libraries (TensorFlow, CNTK, Theano etc.)
  - Python deep learning library
  - Basically an API  “designed for human beings, not machines”
- **TensorFlow** is an open-source deep learning library developed by Google
- When TensorFlow 2.0 was released in late 2019, it adopted Keras as the official API for TensorFlow 2.0
  - Can import Keras directly from TensorFlow library

# Why Keras?

- Easy to use as it provides simple APIs to build deep learning models
- Provide clear & actionable error messages
- Easy to test new ideas, build prototype models, and deploy to production
- Strong multi-GPU & distributed training support
- Broad adoption in the industry and the research community

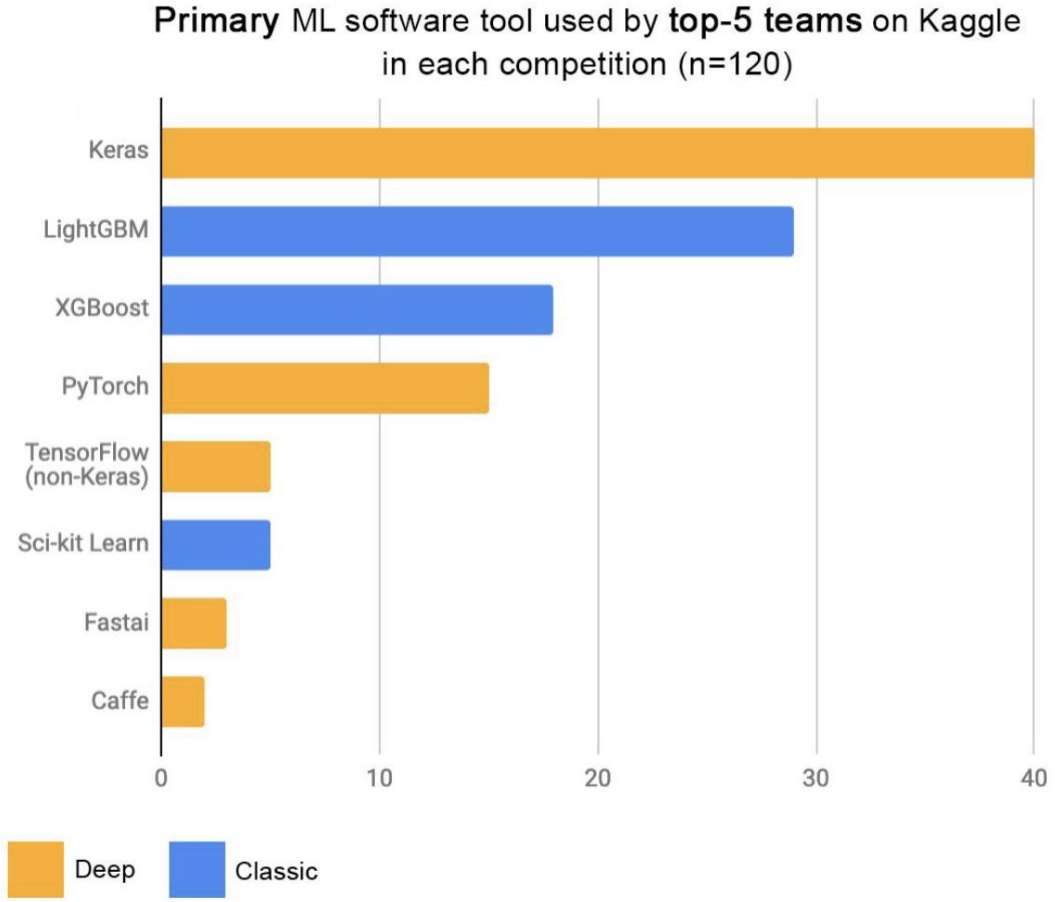


*Year: 2019*

# How to get Keras?

1. Install TensorFlow

```
# Requires the latest pip
$ pip install --upgrade pip

# Current stable release for CPU and GPU
$ pip install tensorflow
```

2. Import Keras

```
import tensorflow as tf
from tensorflow import keras
```

# Let’s use Keras to build a simple ANN model and see how it works

# Let's Connect!

***bit.ly/admond-linkedin***

***bit.ly/admond-medium***

***bit.ly/admond-twitter***

***bit.ly/admond-instagram***

***bit.ly/admond-website***

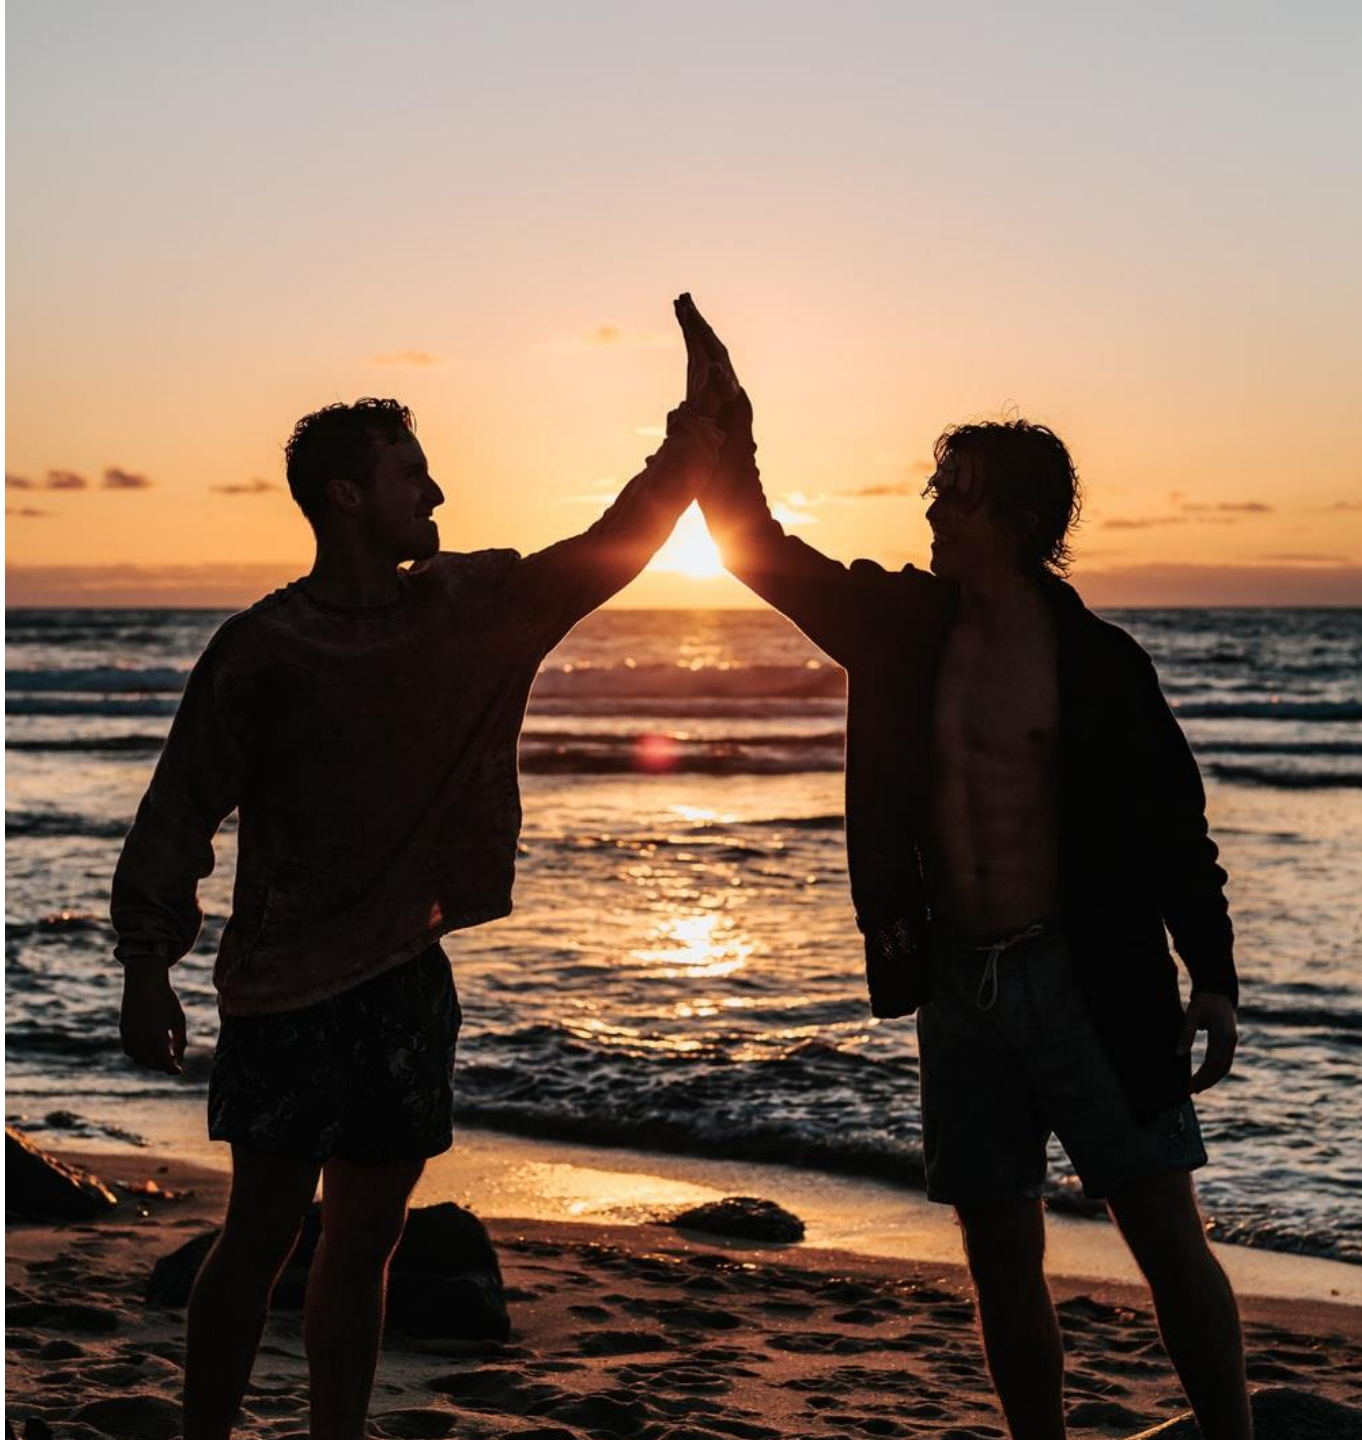

# THANK YOU!

Ask me a question!

Type something....